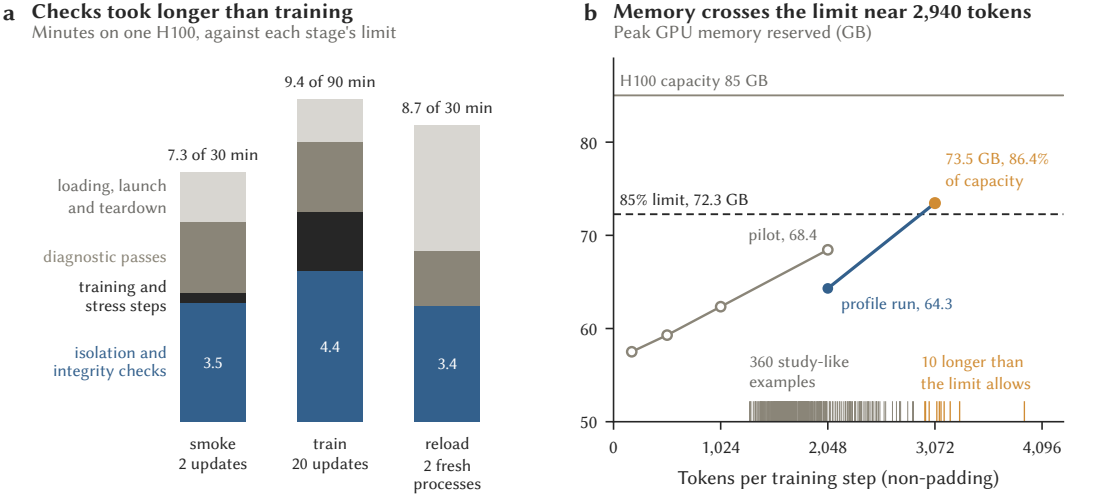
a
Checks took longer than training
Minutes on one H100, against each stage's limit
9.4 of 90 min
8.7 of 30 min
7.3 of 30 min
loading, launch and teardown
diagnostic passes
training and stress steps
isolation and integrity checks
3.5
4.4
3.4
smoke 2 updates
train 20 updates
reload 2 fresh processes
b
Memory crosses the limit near 2,940 tokens
Peak GPU memory reserved (GB)
H100 capacity 85 GB
85% limit, 72.3 GB
73.5 GB, 86.4% of capacity
pilot, 68.4
profile run, 64.3
360 study-like examples
10 longer than the limit allows
80
70
60
50
0
1,024
2,048
3,072
4,096
Tokens per training step (non-padding)