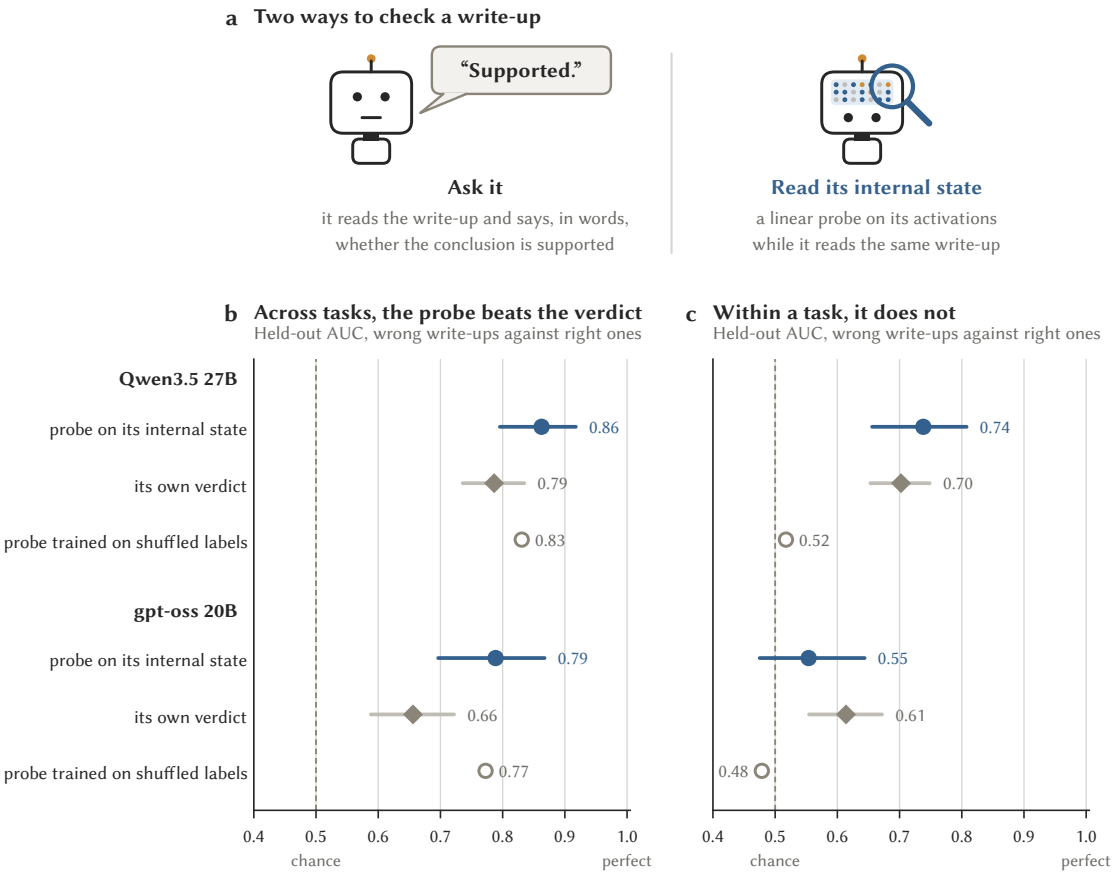
a Two ways to check a write-up
“Supported.”
Ask it
it reads the write-up and says, in words, whether the conclusion is supported
Read its internal state
a linear probe on its activations while it reads the same write-up
b Across tasks, the probe beats the verdict
Held-out AUC, wrong write-ups against right ones
Qwen3.5 27B
probe on its internal state
0.86
its own verdict
0.79
probe trained on shuffled labels
0.83
gpt-oss 20B
probe on its internal state
0.79
its own verdict
0.66
probe trained on shuffled labels
0.77
0.4
0.5
0.6
0.7
0.8
0.9
1.0
chance
perfect
c Within a task, it does not
Held-out AUC, wrong write-ups against right ones
0.74
0.70
0.52
0.55
0.61
0.48
0.4
0.5
0.6
0.7
0.8
0.9
1.0
chance
perfect