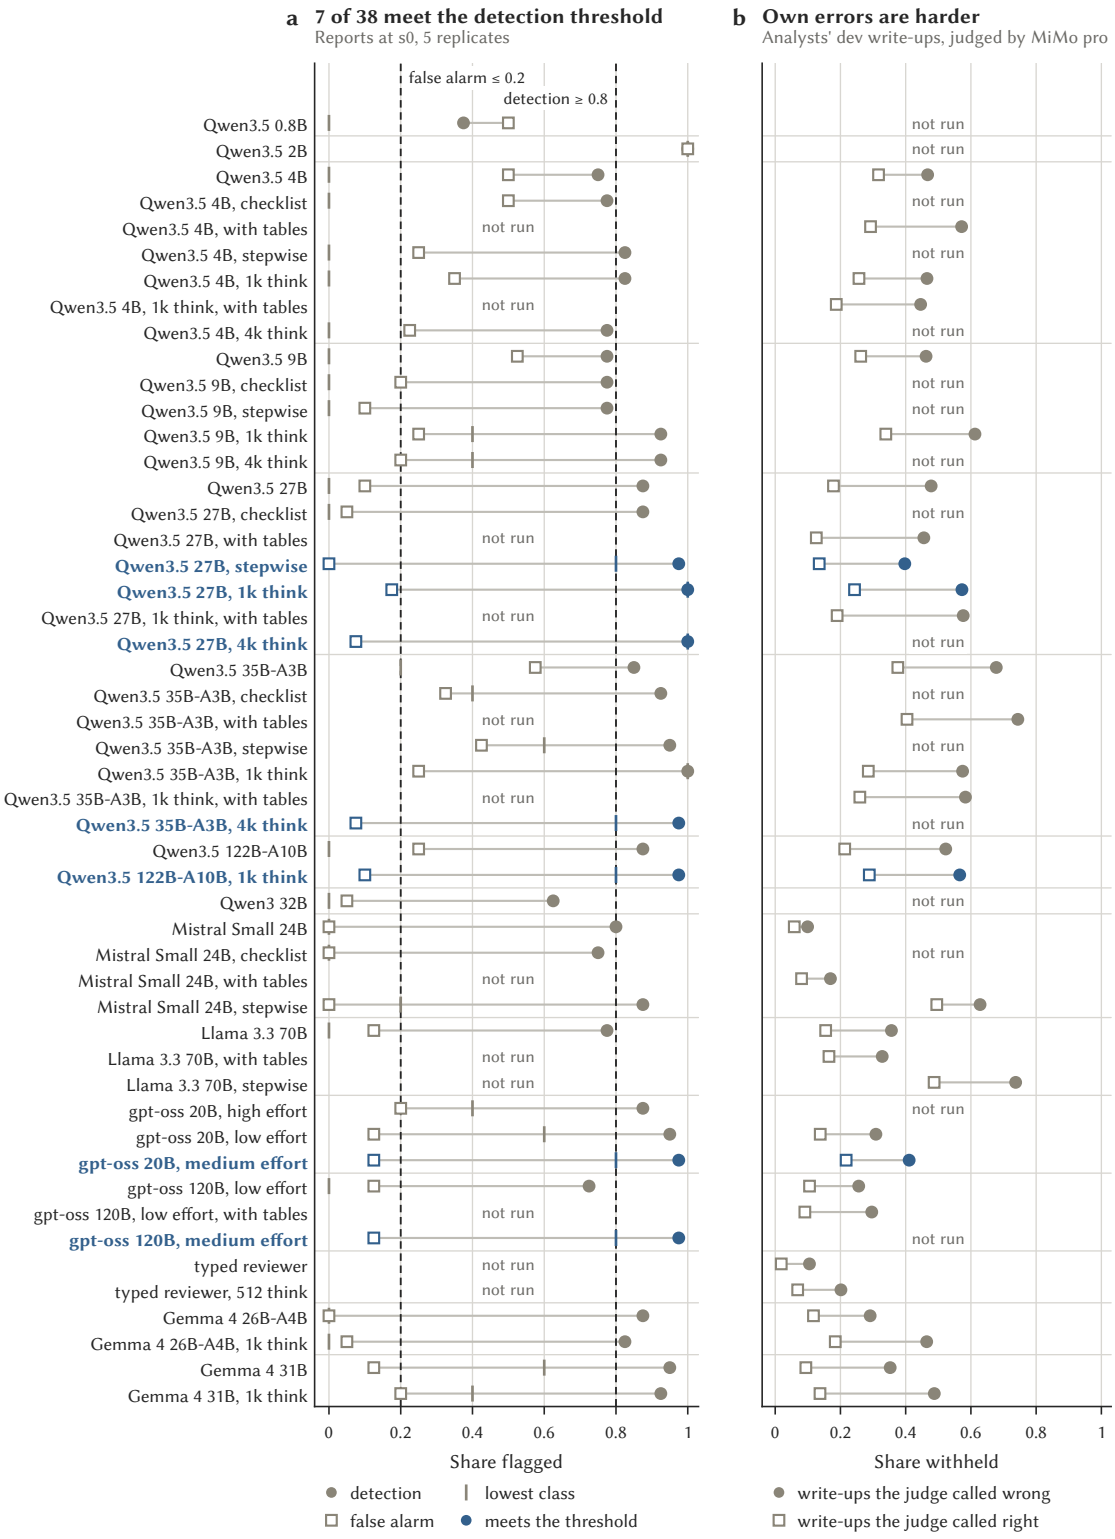
a
7 of 38 meet the detection threshold
Reports at s0, 5 replicates
false alarm ≤ 0.2
detection ≥ 0.8
Qwen3.5 0.8B
Qwen3.5 2B
Qwen3.5 4B
Qwen3.5 4B, checklist
Qwen3.5 4B, with tables
Qwen3.5 4B, stepwise
Qwen3.5 4B, 1k think
Qwen3.5 4B, 1k think, with tables
Qwen3.5 4B, 4k think
Qwen3.5 9B
Qwen3.5 9B, checklist
Qwen3.5 9B, stepwise
Qwen3.5 9B, 1k think
Qwen3.5 9B, 4k think
Qwen3.5 27B
Qwen3.5 27B, checklist
Qwen3.5 27B, with tables
Qwen3.5 27B, stepwise
Qwen3.5 27B, 1k think
Qwen3.5 27B, 1k think, with tables
Qwen3.5 27B, 4k think
Qwen3.5 35B-A3B
Qwen3.5 35B-A3B, checklist
Qwen3.5 35B-A3B, with tables
Qwen3.5 35B-A3B, stepwise
Qwen3.5 35B-A3B, 1k think
Qwen3.5 35B-A3B, 1k think, with tables
Qwen3.5 35B-A3B, 4k think
Qwen3.5 122B-A10B
Qwen3.5 122B-A10B, 1k think
Qwen3 32B
Mistral Small 24B
Mistral Small 24B, checklist
Mistral Small 24B, with tables
Mistral Small 24B, stepwise
Llama 3.3 70B
Llama 3.3 70B, with tables
Llama 3.3 70B, stepwise
gpt-oss 20B, high effort
gpt-oss 20B, low effort
gpt-oss 20B, medium effort
gpt-oss 120B, low effort
gpt-oss 120B, low effort, with tables
gpt-oss 120B, medium effort
typed reviewer
typed reviewer, 512 think
Gemma 4 26B-A4B
Gemma 4 26B-A4B, 1k think
Gemma 4 31B
Gemma 4 31B, 1k think
not run
0
0.2
0.4
0.6
0.8
1
Share flagged
detection
lowest class
false alarm
meets the threshold
b
Own errors are harder
Analysts' dev write-ups, judged by MiMo pro
not run
0
0.2
0.4
0.6
0.8
1
Share withheld
write-ups the judge called wrong
write-ups the judge called right