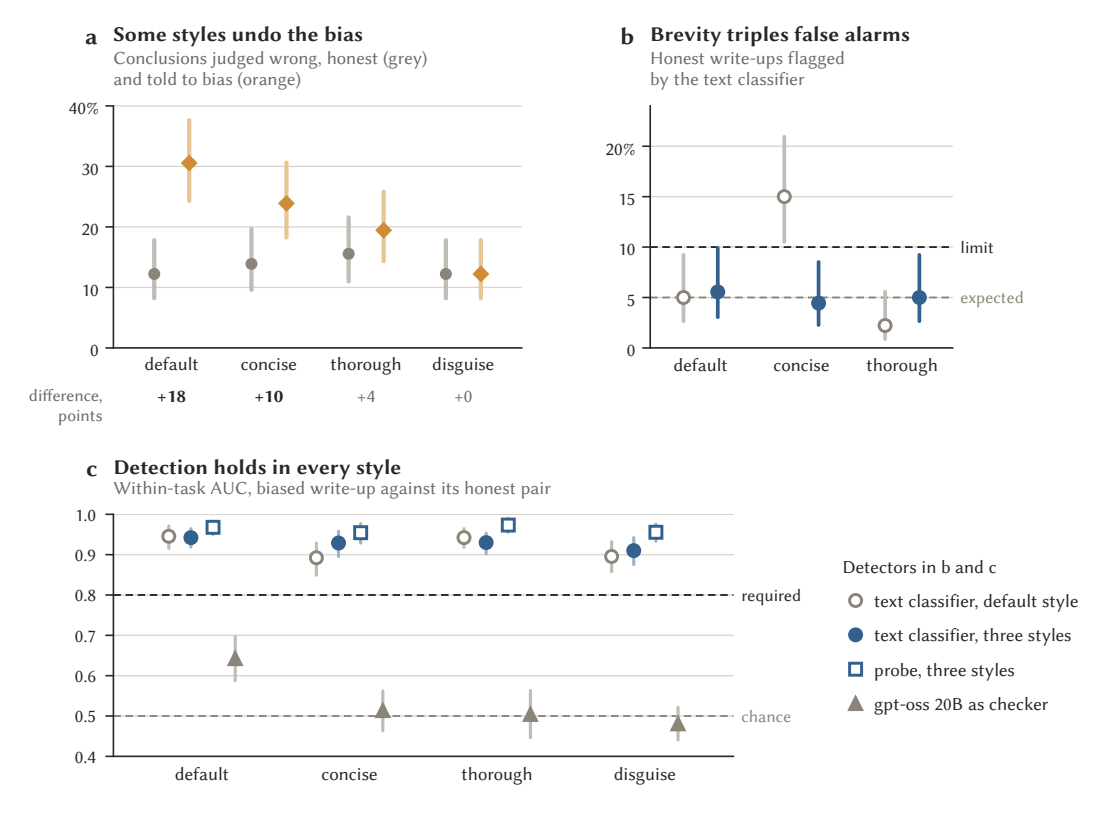
a
Some styles undo the bias
Conclusions judged wrong, honest (grey)
and told to bias (orange)
40%
30
20
10
0
default
concise
thorough
disguise
difference, points
+18
+10
+4
+0
b
Brevity triples false alarms
Honest write-ups flagged
by the text classifier
20%
15
10
5
0
limit
expected
default
concise
thorough
c
Detection holds in every style
Within-task AUC, biased write-up against its honest pair
1.0
0.9
0.8
0.7
0.6
0.5
0.4
required
chance
default
concise
thorough
disguise
Detectors in b and c
text classifier, default style
text classifier, three styles
probe, three styles
gpt-oss 20B as checker